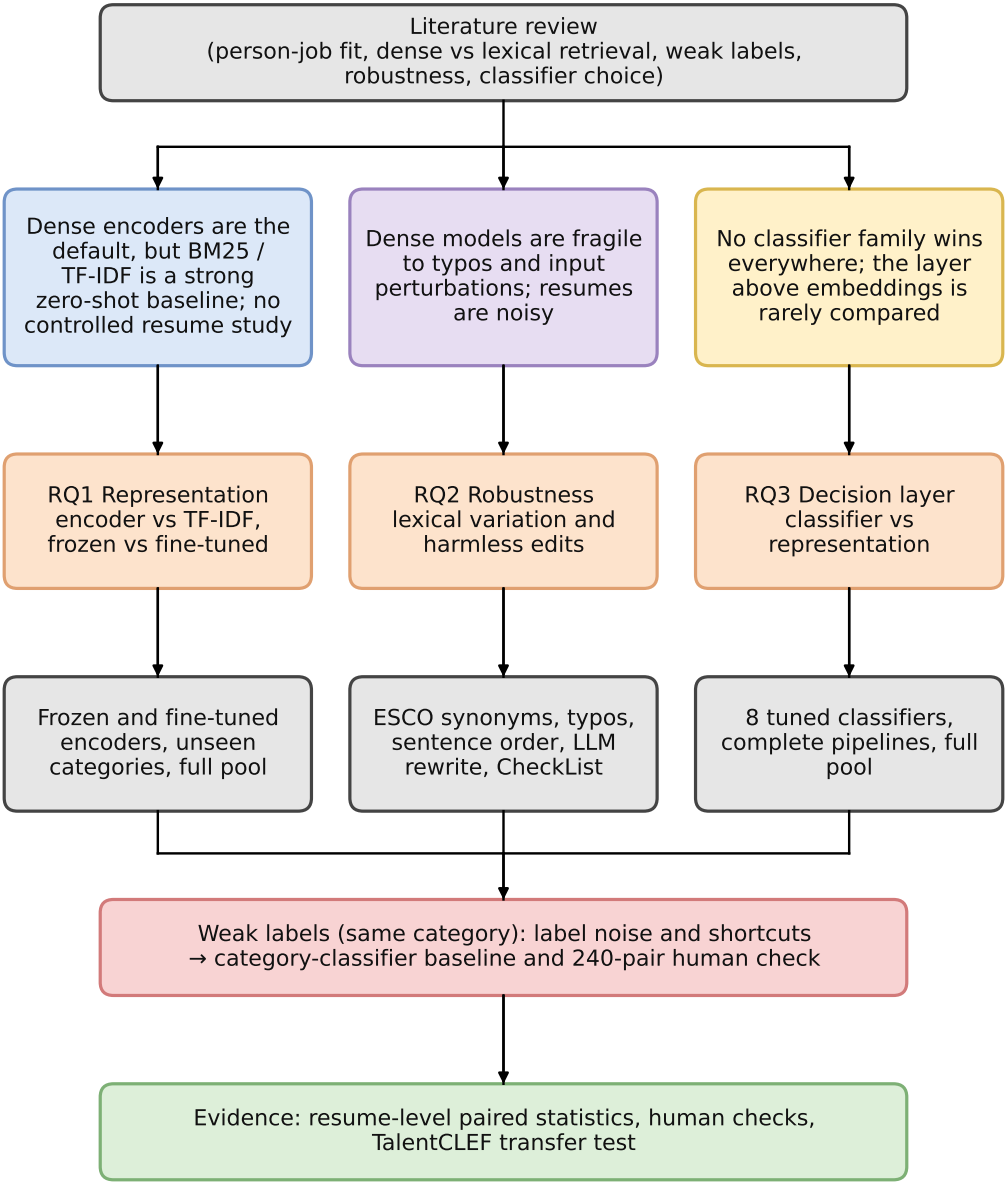
Literature review
(person-job fit, dense vs lexical retrieval, weak labels, robustness, classifier choice)
Dense encoders are the default, but BM25 / TF-IDF is a strong zero-shot baseline; no controlled resume study
Dense models are fragile to typos and input perturbations; resumes are noisy
No classifier family wins everywhere; the layer above embeddings is rarely compared
RQ1 Representation encoder vs TF-IDF, frozen vs fine-tuned
RQ2 Robustness lexical variation and harmless edits
RQ3 Decision layer classifier vs representation
Frozen and fine-tuned encoders, unseen categories, full pool
ESCO synonyms, typos, sentence order, LLM rewrite, CheckList
8 tuned classifiers, complete pipelines, full pool
Weak labels (same category): label noise and shortcuts
→ category-classifier baseline and 240-pair human check
Evidence: resume-level paired statistics, human checks, TalentCLEF transfer test