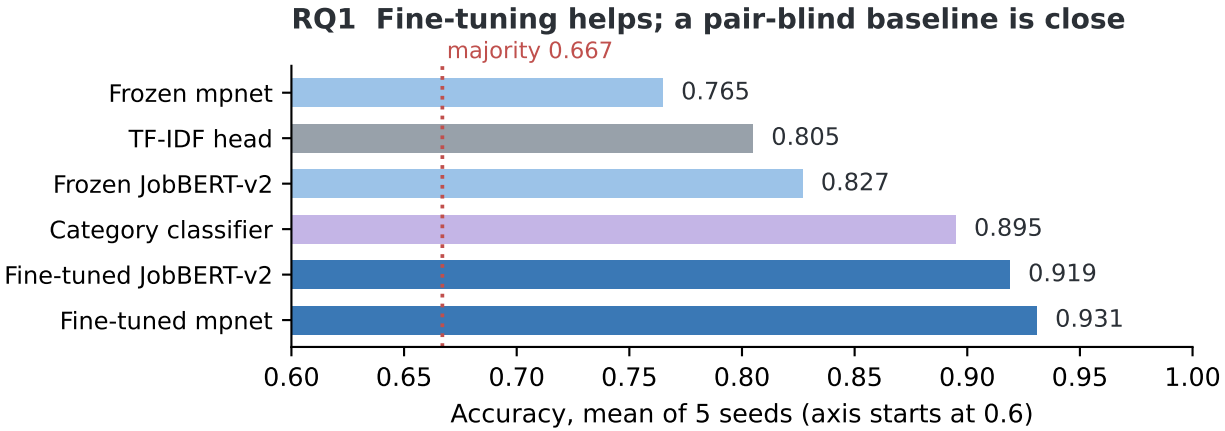
RQ1 Fine-tuning helps; a pair-blind baseline is close
majority 0.667
Frozen mpnet
TF-IDF head
Frozen JobBERT-v2
Category classifier
Fine-tuned JobBERT-v2
Fine-tuned mpnet
0.765
0.805
0.827
0.895
0.919
0.931
0.60
0.65
0.70
0.75
0.80
0.85
0.90
0.95
1.00
Accuracy, mean of 5 seeds (axis starts at 0.6)

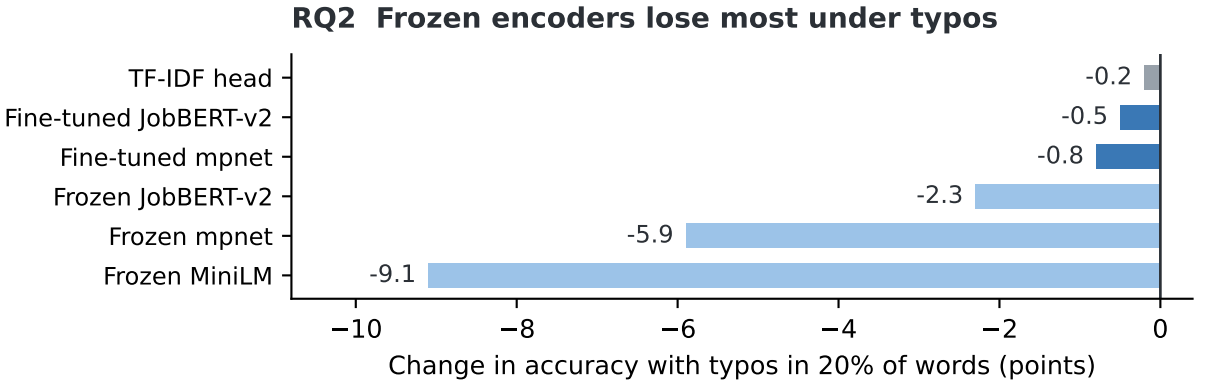
RQ2 Frozen encoders lose most under typos
TF-IDF head
Fine-tuned JobBERT-v2
Fine-tuned mpnet
Frozen JobBERT-v2
Frozen mpnet
Frozen MiniLM
-0.2
-0.5
-0.8
-2.3
-5.9
-9.1
−10
−8
−6
−4
−2
0
Change in accuracy with typos in 20% of words (points)

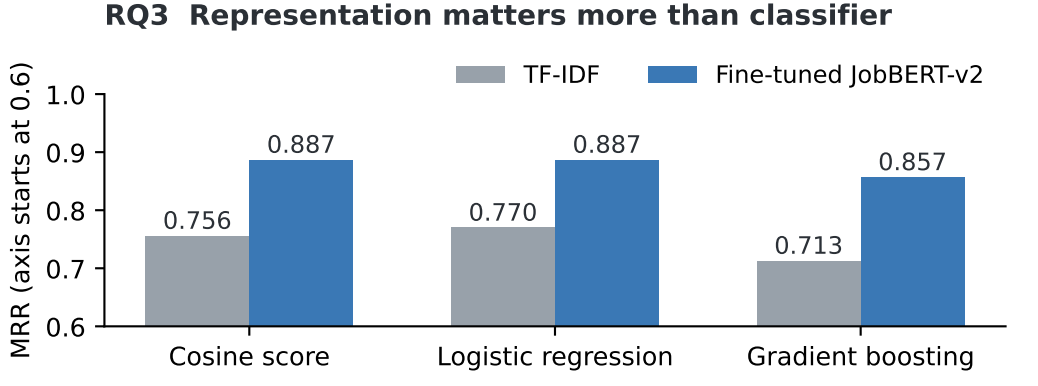
RQ3 Representation matters more than classifier
TF-IDF
Fine-tuned JobBERT-v2
MRR (axis starts at 0.6)
1.0
0.9
0.8
0.7
0.6
0.756
0.887
0.770
0.887
0.713
0.857
Cosine score
Logistic regression
Gradient boosting